

# TERM DEPOSIT PREDICTION

A Machine Learning Project to predict if a client will subscribe (yes/no) to a term deposit.

In this project I will demonstrate how to build a model, to predict which clients will subscribing to a term deposit, with inception of machine learning. In the first part we will deal with the description and visualization of the analysed data, and in the second we will go to data classification models.

## features

- age : age (numeric)
- job : type of job
- marital : marital status
- education : education
- default: has credit in default?
- balance in bank account
- housing: has housing loan?
- loan: has personal loan?
- day\_of\_week: last contact day of the week
- month: last contact month of year
- duration: last contact duration, in seconds
- campaign: number of contacts performed during this campaign and for this client (numeric, includes last contact)
- pdays: number of days that passed by after the client was last contacted from a previous campaign
- previous: number of contacts performed before this campaign and for this client
- y : outcome of the previous marketing y - has the client subscribed a term deposit?

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

df=pd.read_csv('/content/dataset_project.csv (1).csv')
df
```

	age	job	marital	education	default	balance
housing loan \						
0	999.0	management	married	tertiary	no	2143.0
yes no						
1	44.0	technician	single	secondary	no	29.0
yes no						
2	33.0	entrepreneur	married	secondary	no	2.0
yes yes						
3	47.0	blue-collar	married	unknown	no	1506.0
yes no						
4	33.0	unknown	single	unknown	no	1.0
no no						
...	...	...	...	...	...	...
...	...	...	...	...	...	...
45206	51.0	technician	married	tertiary	no	825.0
no no						
45207	71.0	retired	divorced	primary	no	1729.0
no no						
45208	72.0	retired	married	secondary	no	5715.0
no no						
45209	57.0	blue-collar	married	secondary	no	668.0
no no						
45210	37.0	entrepreneur	married	secondary	no	2971.0
no no						

	day	month	duration	campaign	pdays	previous	y
0	5	may	261	1	-1	0	no
1	5	may	151	1	-1	0	no
2	5	may	76	1	-1	0	no
3	5	may	92	1	-1	0	no
4	5	may	198	1	-1	0	no
...	...	...	...	...	...	...	...
45206	17	nov	977	3	-1	0	yes
45207	17	nov	456	2	-1	0	yes
45208	17	nov	1127	5	184	3	yes
45209	17	nov	508	4	-1	0	no
45210	17	nov	361	2	188	11	no

[45211 rows x 15 columns]

df.head()

	age	job	marital	education	default	balance	housing
loan day \							
0	999.0	management	married	tertiary	no	2143.0	yes
no 5							
1	44.0	technician	single	secondary	no	29.0	yes
no 5							
2	33.0	entrepreneur	married	secondary	no	2.0	yes
yes 5							

3	47.0	blue-collar	married	unknown	no	1506.0	yes
no	5						
4	33.0	unknown	single	unknown	no	1.0	no
no	5						

	month	duration	campaign	pdays	previous	y
0	may	261	1	-1	0	no
1	may	151	1	-1	0	no
2	may	76	1	-1	0	no
3	may	92	1	-1	0	no
4	may	198	1	-1	0	no

df.tail()

	age	job	marital	education	default	balance
housing loan \						
45206	51.0	technician	married	tertiary	no	825.0
no	no					
45207	71.0	retired	divorced	primary	no	1729.0
no	no					
45208	72.0	retired	married	secondary	no	5715.0
no	no					
45209	57.0	blue-collar	married	secondary	no	668.0
no	no					
45210	37.0	entrepreneur	married	secondary	no	2971.0
no	no					

	day	month	duration	campaign	pdays	previous	y
45206	17	nov	977	3	-1	0	yes
45207	17	nov	456	2	-1	0	yes
45208	17	nov	1127	5	184	3	yes
45209	17	nov	508	4	-1	0	no
45210	17	nov	361	2	188	11	no

df.shape

(45211, 15)

df.dtypes

age	float64
job	object
marital	object
education	object
default	object
balance	float64
housing	object
loan	object
day	int64
month	object
duration	int64

```
campaign      int64
pdays        int64
previous       int64
y             object
dtype: object
```

```
df.describe()
```

	age	balance	day	duration
campaign \				
count	45202.000000	45208.000000	45211.000000	45211.000000
mean	40.954714	1362.346620	15.806419	258.163080
std	2.763841	3044.852387	8.322476	257.527812
min	-1.000000	-8019.000000	1.000000	0.000000
25%	33.000000	72.000000	8.000000	103.000000
50%	39.000000	448.000000	16.000000	180.000000
75%	48.000000	1428.000000	21.000000	319.000000
max	999.000000	102127.000000	31.000000	4918.000000

	pdays	previous
count	45211.000000	45211.000000
mean	40.197828	0.580323
std	100.128746	2.303441
min	-1.000000	0.000000
25%	-1.000000	0.000000
50%	-1.000000	0.000000
75%	-1.000000	0.000000
max	871.000000	275.000000

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         45202 non-null  float64
1   job         45211 non-null  object
2   marital     45211 non-null  object
3   education   45211 non-null  object
4   default     45211 non-null  object
5   balance     45208 non-null  float64
```

```
6   housing      45211 non-null object
7   loan         45211 non-null object
8   day          45211 non-null int64
9   month        45211 non-null object
10  duration     45211 non-null int64
11  campaign     45211 non-null int64
12  pdays        45211 non-null int64
13  previous     45211 non-null int64
14  y            45211 non-null object
dtypes: float64(2), int64(5), object(8)
memory usage: 5.2+ MB
```

```
df.isna().sum()
```

```
age      9
job       0
marital   0
education 0
default   0
balance   3
housing   0
loan       0
day        0
month      0
duration   0
campaign   0
pdays     0
previous   0
y          0
dtype: int64
```

```
for col in df.columns: #To Know unique values
    print(col, ' ', df[col].nunique())
```

```
age      79
job      12
marital    3
education  4
default    2
balance  7168
housing    2
loan       2
day       31
month     12
duration  1573
campaign   48
pdays    559
previous   41
y         2
```

```
df.columns
Index(['age', 'job', 'marital', 'education', 'default', 'balance',
       'housing',
       'loan', 'day', 'month', 'duration', 'campaign', 'pdays',
       'previous',
       'y'],
      dtype='object')
```

```
for col in df.columns:
    print(col)
    print(df[col].nunique())
    print('-'*100)
```

```
age
79
```

```
-----
```

```
job
12
```

```
-----
```

```
marital
3
```

```
-----
```

```
education
4
```

```
-----
```

```
default
2
```

```
-----
```

```
balance
7168
```

```
-----
```

```
housing
2
```

```
-----
```

```
loan
2
```

```
-----
```

```
day
31
```

```
-----
```

month  
12

duration  
1573

campaign  
48

pdays  
559

previous  
41

y  
2

```
for col in df.columns:
    print(col)
    print(df[col].unique())
    print('-'*100)
```

age  
[999. 44. 33. 47. 35. 28. nan 58. 43. 41. 29. 53. 57. 51.  
45. 60. 56. 32. 25. 40. 39. 52. 46. 36. 49. 59. 37. 50.  
54. 55. 48. 31. 42. 30. 27. 34. 38. 23. 26. 61. 22. 24.  
21. 20. 66. 62. 83. 75. 67. 70. 65. 68. 64. 69. 72. 71.  
19. 76. 85. 63. 90. 82. 73. 74. 78. 80. 94. 79. 77. 86.  
95. 81. 18. 89. 84. 87. 92. 93. 88. -1.]

job  
['management' 'technician' 'entrepreneur' 'blue-collar' 'unknown'  
'retired' 'admin.' 'services' 'self-employed' 'unemployed'  
'housemaid'  
'student']

marital  
['married' 'single' 'divorced']

```
education
['tertiary' 'secondary' 'unknown' 'primary']
-----

default
['no' 'yes']
-----

balance
[2.1430e+03 2.9000e+01 2.0000e+00 ... 8.2050e+03 1.4204e+04
1.6353e+04]
-----

housing
['yes' 'no']
-----

loan
['no' 'yes']
-----

day
[ 5  6  7  8  9 12 13 14 15 16 19 20 21 23 26 27 28 29 30  2  3  4 11
17
 18 24 25  1 10 22 31]
-----

month
['may' 'jun' 'jul' 'aug' 'oct' 'nov' 'dec' 'jan' 'feb' 'mar' 'apr'
'sep']
-----

duration
[ 261  151   76 ... 1298 1246 1556]
-----

campaign
[ 1  2  3  5  4  6  7  8  9 10 11 12 13 19 14 24 16 32 18 22 15 17 25
21
 43 51 63 41 26 28 55 50 38 23 20 29 31 37 30 46 27 58 33 35 34 36 39
44]
-----

pdays
[ -1 151 166  91  86 143 147  89 140 176 101 174 170 167 195 165 129
188
 196 172 118 119 104 171 117 164 132 131 123 159 186 111 115 116 173
178
 110 152  96 103 150 175 193 181 185 154 145 138 126 180 109 158 168
```



97  
182 127 130 194 125 105 102 26 179 28 183 155 112 120 137 124 187  
190  
113 162 134 169 189 8 144 191 184 177 5 99 133 93 92 10 100  
156  
198 106 153 146 128 7 121 160 107 90 27 197 136 139 122 157 149  
135  
30 114 98 192 163 34 95 141 31 199 94 108 29 268 247 253 226  
244  
239 245 204 231 238 258 230 254 265 71 223 246 250 266 240 205 261  
259  
241 260 234 251 225 161 237 262 248 255 220 227 206 224 249 235 228  
263  
2 270 232 252 207 200 269 233 256 273 272 242 264 208 214 222 271  
203  
221 202 216 201 257 229 210 217 75 213 73 76 267 211 215 77 236  
82  
6 209 274 1 243 212 275 80 276 9 279 12 280 88 277 85 84  
219  
24 21 282 41 294 49 329 307 303 331 308 300 64 314 287 330 332  
302  
323 318 333 60 326 335 313 312 305 325 327 336 309 328 322 39 316  
292  
295 310 306 320 317 289 57 321 142 339 301 315 337 334 340 319 17  
74  
148 341 299 344 342 324 345 346 304 281 343 338 14 347 15 291 348  
349  
285 350 284 25 283 278 81 4 87 83 79 70 13 293 37 78 63  
22  
296 355 66 19 35 360 357 354 351 362 358 365 298 286 364 363 47  
361  
288 366 356 352 359 297 367 353 368 42 290 67 371 370 369 50 36  
373  
374 372 311 375 378 59 379 40 18 43 20 69 38 385 56 55 44  
391  
72 390 32 62 399 393 65 377 395 388 389 386 61 412 405 434 394  
382  
459 440 397 383 68 461 462 463 422 51 457 430 442 403 454 428 392  
410  
401 474 475 477 478 54 476 380 479 45 46 495 58 48 518 52 515  
520  
511 536 387 218 33 544 435 436 555 433 446 558 469 616 561 553 384  
592  
467 585 480 421 667 626 426 595 381 376 648 521 452 449 633 398 53  
460  
670 551 414 557 687 404 651 686 425 504 578 674 416 586 411 756 450  
745  
514 417 424 776 396 683 529 439 415 456 407 458 532 481 791 701 531  
792

```

413 445 535 784 419 455 491 431 542 470 472 717 437 3 782 728 828
524
562 761 492 775 579 493 464 760 466 465 656 831 490 432 655 427 749
838
769 587 778 854 779 850 771 594 842 589 603 484 489 486 409 444 680
808
485 503 690 772 774 526 420 528 500 826 804 508 547 805 541 543 871
550
530]
-----
-----

```

previous

```

[ 0 3 1 4 2 11 16 6 5 10 12 7 18 9 21 8 14
15
26 37 13 25 20 27 17 23 38 29 24 51 275 22 19 30 58
28
32 40 55 35 41]
-----
-----

```

y  
['no' 'yes']

## FEATURE SELECTION

```
df.corr()
```

```

-----
-----
NameError                                Traceback (most recent call
last)
<ipython-input-1-2f6f6606aa2c> in <cell line: 1>()
----> 1 df.corr()

```

NameError: name 'df' is not defined

```

plt.figure(figsize=(10,5))
sns.heatmap(df.corr(),annot=True)

```

```

<ipython-input-15-fe2252758f39>:2: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it
will default to False. Select only valid columns or specify the value
of numeric_only to silence this warning.

```

```
    sns.heatmap(df.corr(),annot=True)
```

```
<Axes: >
```



45208	55	5	1	1	0	5455	0	0
16								
45209	40	1	1	1	0	1584	0	0
16								
45210	20	2	1	1	0	3779	0	0
16								

	month	duration	campaign	pdays	previous	y
0	8	261	0	0	0	0
1	8	151	0	0	0	0
2	8	76	0	0	0	0
3	8	92	0	0	0	0
4	8	198	0	0	0	0
...	...	...	...	...	...	..
45206	9	975	2	0	0	1
45207	9	456	1	0	0	1
45208	9	1116	4	181	3	1
45209	9	508	3	0	0	0
45210	9	361	1	185	11	0

[45211 rows x 15 columns]

```
x_discrete=df_copy[['job', 'marital', 'education', 'default',
'housing', 'loan', 'month']]
x_continuous=df_copy[['age', 'balance', 'day', 'duration',
'campaign', 'pdays', 'previous']]
y_copy=df_copy.iloc[:, -1]
```

*#Chi Square Test for feature selection in case of discrete input labels*

```
from sklearn.feature_selection import chi2
score1=chi2(x_discrete,y_copy)
score1
```

```
(array([182.45226044, 29.76606652, 90.61772256, 22.31387496,
388.94971474, 176.51613693, 44.32190507]),
array([1.41257633e-41, 4.87449759e-08, 1.74292216e-21, 2.31527677e-
06,
1.40128480e-86, 2.79337524e-40, 2.78581489e-11]))
```

```
f_valuel=pd.Series(score1[0],index=x_discrete.columns)
f_valuel.sort_values(ascending=False)
```

```
housing      388.949715
job          182.452260
loan         176.516137
education    90.617723
month        44.321905
marital      29.766067
default      22.313875
dtype: float64
```

```
p_value1=pd.Series(score1[1],index=x_discrete.columns)
p_value1.sort_values(ascending=False)
```

```
default      2.315277e-06
marital      4.874498e-08
month        2.785815e-11
education    1.742922e-21
loan         2.793375e-40
job          1.412576e-41
housing      1.401285e-86
dtype: float64
```

```
#Anova Test for feature selection in case of continuous input labels
from sklearn.feature_selection import f_classif
score2=f_classif(x_continuous,y_copy)
score2
```

```
(array([ 26.58010222, 311.62865296,  36.35900965, 9021.94787625,
        245.7274607 , 470.12762293, 599.58462133]),
 array([2.53897219e-007, 1.65228194e-069, 1.65388016e-009,
        0.00000000e+000,
        3.10161380e-055, 1.01716123e-103, 1.48875810e-131]))
```

```
f_value2=pd.Series(score2[0],index=x_continuous.columns)
f_value2.sort_values(ascending=False)
```

```
duration      9021.947876
previous      599.584621
pdays        470.127623
balance       311.628653
campaign      245.727461
day           36.359010
age           26.580102
dtype: float64
```

```
p_value2=pd.Series(score2[1],index=x_continuous.columns)
p_value2.sort_values(ascending=False)
```

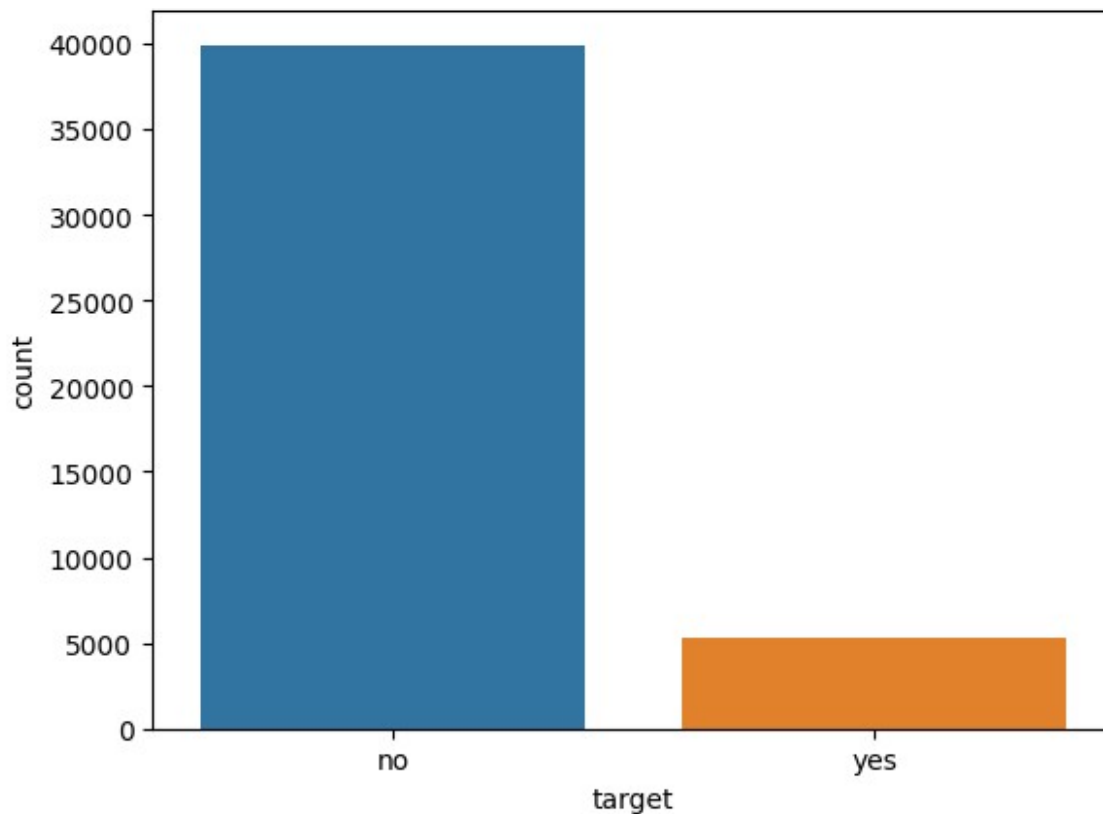
```
age          2.538972e-07
day          1.653880e-09
campaign     3.101614e-55
balance      1.652282e-69
pdays       1.017161e-103
previous     1.488758e-131
duration     0.000000e+00
dtype: float64
```

# VISUALIZATION

```
sns.countplot(x=df['y'])  
plt.xlabel('target')
```



```
df['y'].value_counts()
```



```
no      39922  
yes      5289  
Name: y, dtype: int64
```





Based on this analysis, we can conclude that the dataset is imbalanced.



```
#univariate analysis  
import seaborn as sns  
categorical_cols= [col for col in df.columns if df[col].dtype ==  
'object']  
for i in categorical_cols:  
    plt.figure(figsize=(10,5))  
    sns.countplot(x=i,data=df)  
    plt.title(i)  
    plt.xticks(rotation=45)
```



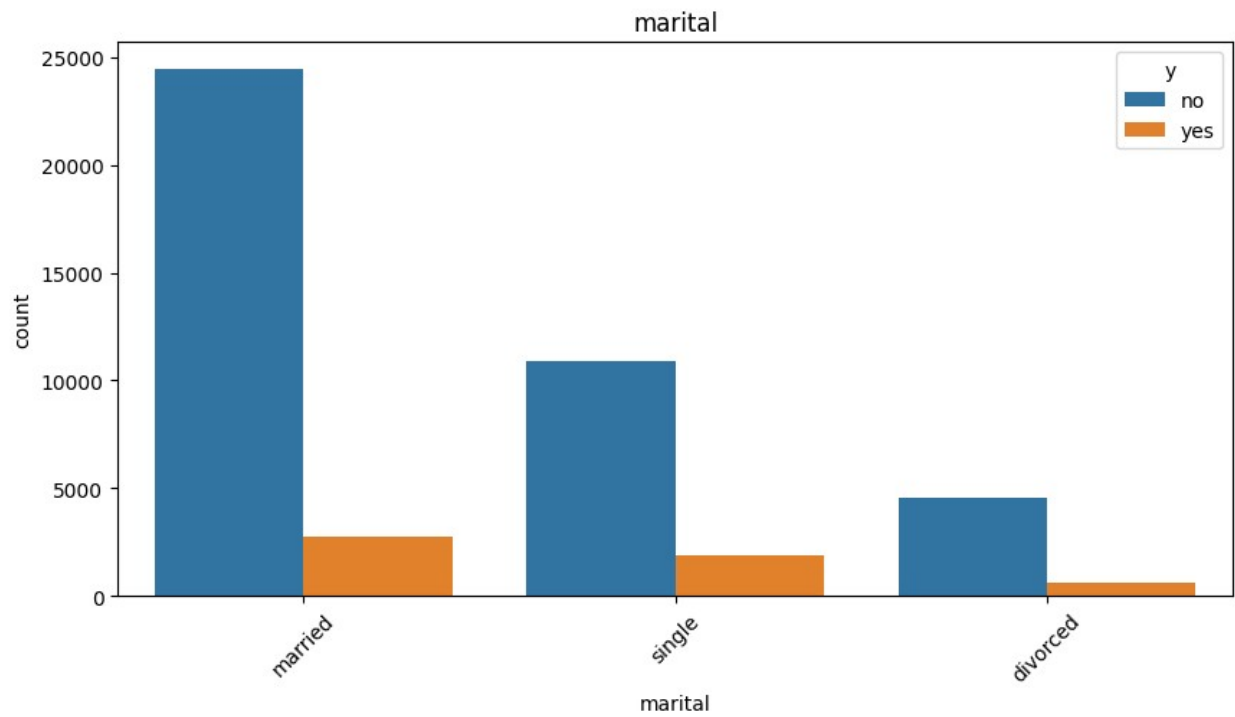







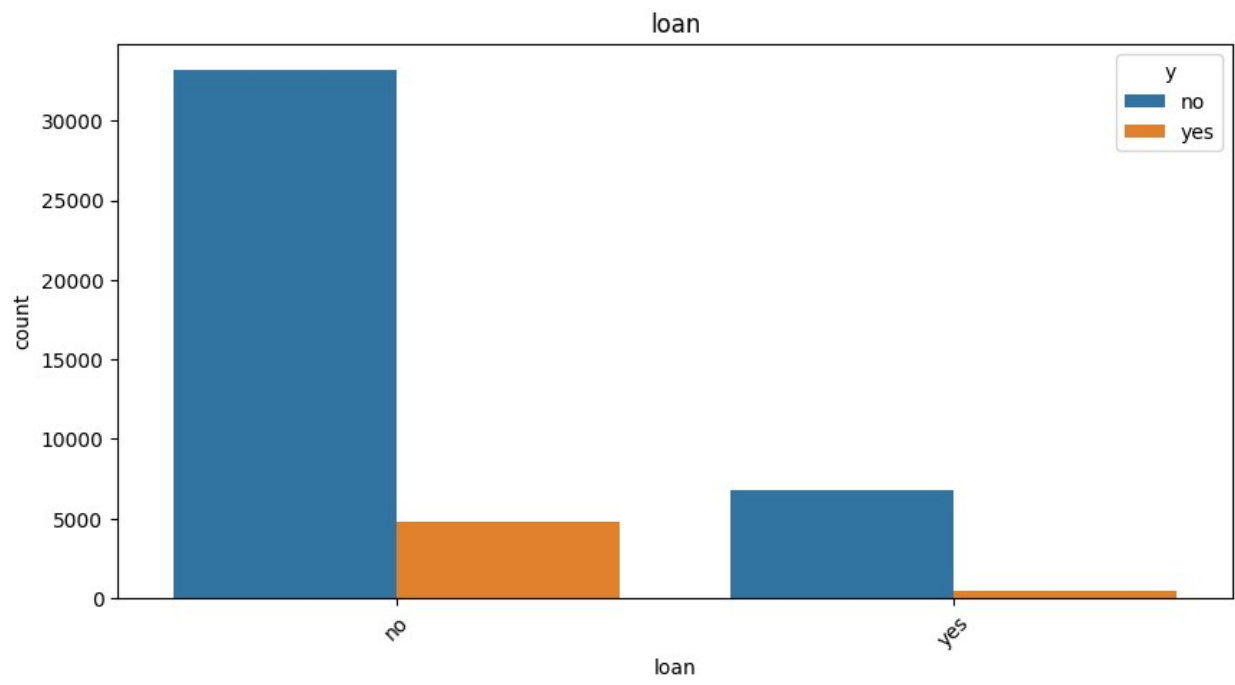







```
#bivariate analysis
for i in categorical_cols:
    plt.figure(figsize=(10,5))
    sns.countplot(x=df[i],hue=df['y'])
    plt.title(i)
    plt.xticks(rotation=45)
```









```
import pandas as pd
contingency_table = pd.crosstab(df['marital'], df['housing'],
                                margins=True, margins_name="Total")
contingency_table
```

housing	no	yes	Total
marital			

divorced	2300	2907	5207
married	11893	15321	27214
single	5888	6902	12790
Total	20081	25130	45211

## HANDLING MISSING VALUES

```
df.isna().sum()
```

```
age          9
job          0
marital      0
education    0
default      0
balance      3
housing      0
loan         0
day          0
month        0
duration     0
campaign     0
pdays       0
previous     0
y            0
dtype: int64
```

```
df['age'] = df['age'].fillna(df['age'].median())    # Filing the null
values with median
df['balance'] = df['balance'].fillna(df['balance'].median()) # Filing
the null values with median
```

In some columns, missing values has been represented as unknown


```
for i in categorical_cols:
    mode=df[i].mode()[0]
    df[i]=df[i].replace('unknown',mode)
```

## DEALING WITH OUTLIERS

```
numerical_cols=['age','balance','day','duration','campaign',
'pdays','previous']
for i in numerical_cols:
    plt.figure(figsize=(8, 6))
    sns.boxplot(x=i , data=df, orient='v', width=0.3)
```


```
plt.title(i)
plt.show()
```

```
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1599:
UserWarning: Vertical orientation ignored with only `x` specified.
warnings.warn(single_var_warning.format("Vertical", "x"))
```




```
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1599:
UserWarning: Vertical orientation ignored with only `x` specified.
warnings.warn(single_var_warning.format("Vertical", "x"))
```







```
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1599:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```



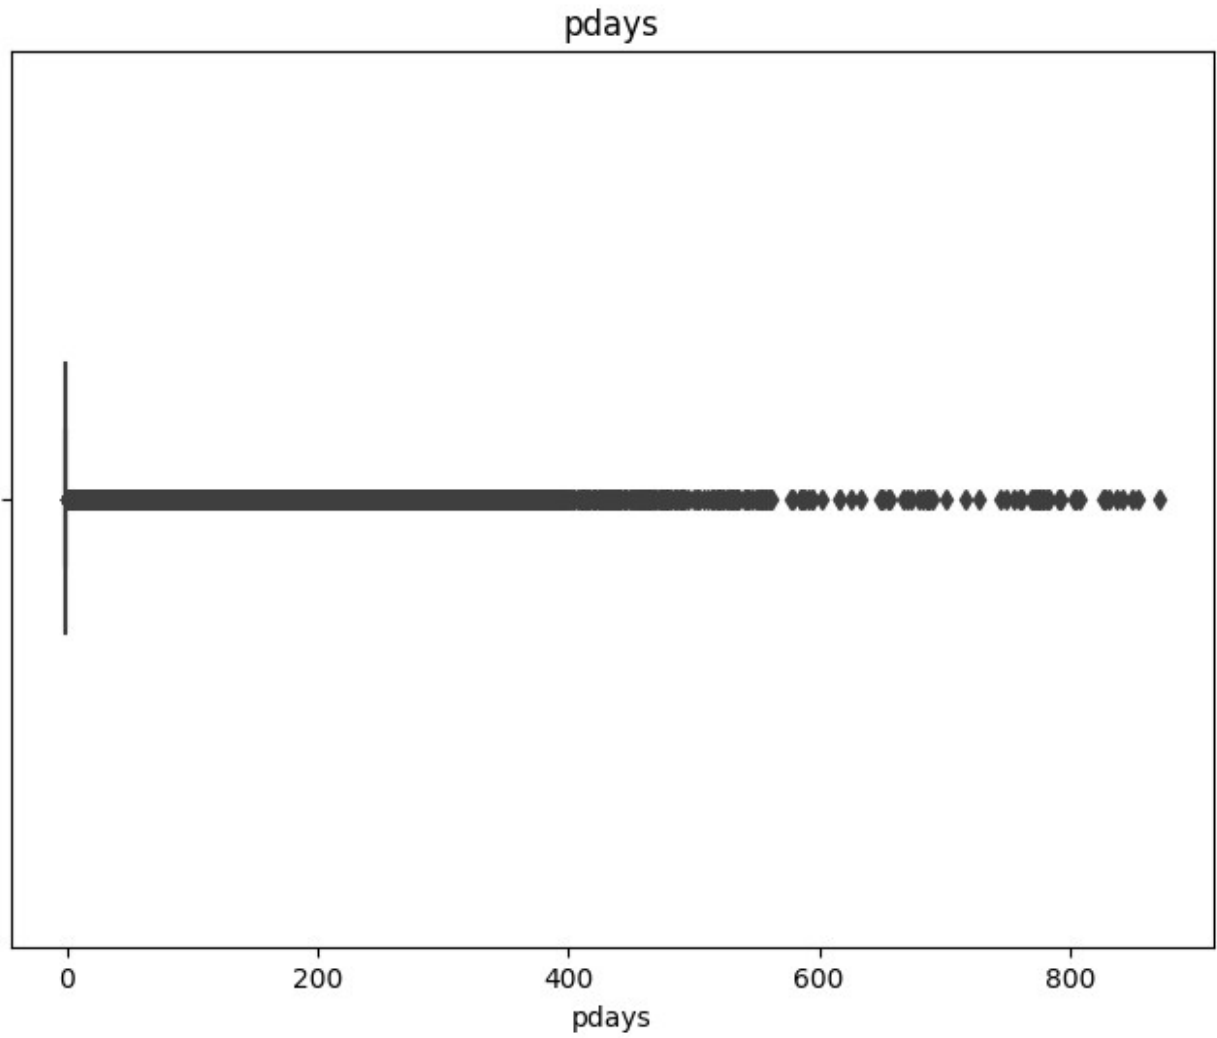
```
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1599:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```




```
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1599:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```



```
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1599:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```



```
/usr/local/lib/python3.10/dist-packages/seaborn/_oldcore.py:1599:  
UserWarning: Vertical orientation ignored with only `x` specified.  
warnings.warn(single_var_warning.format("Vertical", "x"))
```



```
Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
IQR = Q3 - Q1
threshold = 1.5
df = df[~((df < (Q1 - threshold * IQR)) | (df > (Q3 + threshold * IQR))).any(axis=1)]
```

<ipython-input-32-6865bb37bd72>:1: FutureWarning: The default value of numeric\_only in DataFrame.quantile is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
Q1 = df.quantile(0.25)
```

<ipython-input-32-6865bb37bd72>:2: FutureWarning: The default value of numeric\_only in DataFrame.quantile is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
Q3 = df.quantile(0.75)
```

<ipython-input-32-6865bb37bd72>:5: FutureWarning: Automatic reindexing

on DataFrame vs Series comparisons is deprecated and will raise ValueError in a future version. Do `left, right = left.align(right, axis=1, copy=False)` before e.g. `left == right`

```
df= df[~((df < (Q1 - threshold * IQR)) | (df > (Q3 + threshold * IQR))).any(axis=1)]
```

## ENCODING USING GET\_DUMMIES

```
df1=pd.get_dummies(df[['job', 'marital', 'education', 'default', 'housing', 'loan']],drop_first=True)
df1
```

	job_blue-collar	job_entrepreneur	job_housemaid
job_management \			
1	0	0	0
0			
2	0	1	0
0			
3	1	0	0
0			
4	1	0	0
0			
5	0	0	0
1			
...	...	...	...
45196	0	0	0
0			
45197	0	0	0
1			
45198	0	0	0
1			
45202	0	0	0
0			
45209	1	0	0
0			

	job_retired	job_self-employed	job_services	job_student
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0
5	0	0	0	0
...	...	...	...	...
45196	0	0	0	1
45197	0	0	0	0
45198	0	0	0	0
45202	0	0	0	0

45209	0	0	0	0
	job_technician	job_unemployed	marital_married	marital_single
\				
1	1	0	0	1
2	0	0	1	0
3	0	0	1	0
4	0	0	0	1
5	0	0	1	0
...	...	...	...	...
45196	0	0	0	1
45197	0	0	0	1
45198	0	0	1	0
45202	0	0	0	1
45209	0	0	1	0
	education_secondary	education_tertiary	default_yes	
housing_yes \				
1	1	0	0	
1				
2	1	0	0	
1				
3	1	0	0	
1				
4	1	0	0	
0				
5	0	1	0	
1				
...	...	...	...	.
..				
45196	1	0	0	
0				
45197	1	0	0	
1				
45198	0	1	0	
0				
45202	1	0	0	
0				
45209	1	0	0	



0

```
      loan_yes
1          0
2          1
3          0
4          0
5          0
...      ...
45196      0
45197      0
45198      0
45202      0
45209      0
```

[28190 rows x 17 columns]

```
newdf=pd.concat([df,df1],axis=1)
newdf
```

	age	job	marital	education	default	balance	housing
loan \							
1	44.0	technician	single	secondary	no	29.0	yes
no							
2	33.0	entrepreneur	married	secondary	no	2.0	yes
yes							
3	47.0	blue-collar	married	secondary	no	1506.0	yes
no							
4	33.0	blue-collar	single	secondary	no	1.0	no
no							
5	35.0	management	married	tertiary	no	231.0	yes
no							
...	...	...	...	...	...	...	...
...							
45196	25.0	student	single	secondary	no	358.0	no
no							
45197	36.0	management	single	secondary	no	1511.0	yes
no							
45198	37.0	management	married	tertiary	no	1428.0	no
no							
45202	34.0	admin.	single	secondary	no	557.0	no
no							
45209	57.0	blue-collar	married	secondary	no	668.0	no
no							

	day	month	...	job_student	job_technician	job_unemployed	\
1	5	may	...	0	1	0	
2	5	may	...	0	0	0	
3	5	may	...	0	0	0	
4	5	may	...	0	0	0	

5	5	may	...	0	0	0
...	...	...	...	...	...	...
45196	16	nov	...	1	0	0
45197	16	nov	...	0	0	0
45198	16	nov	...	0	0	0
45202	17	nov	...	0	0	0
45209	17	nov	...	0	0	0

	marital_married	marital_single	education_secondary	\
1	0	1	1	
2	1	0	1	
3	1	0	1	
4	0	1	1	
5	1	0	0	
...	...	...	...	
45196	0	1	1	
45197	0	1	1	
45198	1	0	0	
45202	0	1	1	
45209	1	0	1	

	education_tertiary	default_yes	housing_yes	loan_yes
1	0	0	1	0
2	0	0	1	1
3	0	0	1	0
4	0	0	0	0
5	1	0	1	0
...	...	...	...	...
45196	0	0	0	0
45197	0	0	1	0
45198	1	0	0	0
45202	0	0	0	0
45209	0	0	0	0

[28190 rows x 32 columns]

```
newdf1=newdf.drop(['job', 'marital', 'education', 'default',
'housing', 'loan'],axis=1)
newdf1
```

	age	balance	day	month	duration	campaign	pdays	previous
y \								
1	44.0	29.0	5	may	151	1	-1	0
no								
2	33.0	2.0	5	may	76	1	-1	0
no								
3	47.0	1506.0	5	may	92	1	-1	0
no								
4	33.0	1.0	5	may	198	1	-1	0
no								

5	35.0	231.0	5	may	139	1	-1	0
no								
...	...	...	...	...	...	...	...	...
...								
45196	25.0	358.0	16	nov	330	1	-1	0
yes								
45197	36.0	1511.0	16	nov	270	1	-1	0
yes								
45198	37.0	1428.0	16	nov	333	2	-1	0
no								
45202	34.0	557.0	17	nov	224	1	-1	0
yes								
45209	57.0	668.0	17	nov	508	4	-1	0
no								

	job_blue-collar	...	job_student	job_technician
job_unemployed \				
1	0	...	0	1
0				
2	0	...	0	0
0				
3	1	...	0	0
0				
4	1	...	0	0
0				
5	0	...	0	0
0				
...	...	...	...	...
..				
45196	0	...	1	0
0				
45197	0	...	0	0
0				
45198	0	...	0	0
0				
45202	0	...	0	0
0				
45209	1	...	0	0
0				

	marital_married	marital_single	education_secondary	\
1	0	1	1	
2	1	0	1	
3	1	0	1	
4	0	1	1	
5	1	0	0	
...	...	...	...	
45196	0	1	1	
45197	0	1	1	

45198	1	0	0
45202	0	1	1
45209	1	0	1

	education_tertiary	default_yes	housing_yes	loan_yes
1	0	0	1	0
2	0	0	1	1
3	0	0	1	0
4	0	0	0	0
5	1	0	1	0
...	...	...	...	...
45196	0	0	0	0
45197	0	0	1	0
45198	1	0	0	0
45202	0	0	0	0
45209	0	0	0	0

[28190 rows x 26 columns]

```
newdf1['month'] =
df['month'].map({'jan':1,'feb':2,'mar':3,'apr':4,'may':5,'jun':6,'jul':
:7,'aug':8,'sep':9,'oct':10,'nov':11,'dec':12})
#manual encoding for months to keep months as per numerical order
```

## SEPARATING INDEPENDENT AND DEPENDENT VARIABLES


```
x = newdf1.drop('y', axis=1)
y = newdf1['y']
```

## HANDLING IMBALANCED DATASET

```
sns.countplot(x=df['y'])
plt.xlabel('target')

df['y'].value_counts()

no      26593
yes      1597
Name: y, dtype: int64
```



```
from imblearn.over_sampling import SMOTE

smote = SMOTE(sampling_strategy='auto', random_state=42)
x_balanced, y_balanced = smote.fit_resample(x, y)
y_balanced.value_counts()

no      26593
yes     26593
Name: y, dtype: int64
```

## SEPARATING TRAINING DATA & TESTING DATA

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x_balanced,y_balanced,t
est size=0.30,random state=42)
```

# NORMALIZATION

```
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
scaler.fit(x_train)
x_train=scaler.transform(x_train)
x_test=scaler.transform(x_test)
```

# MODEL CREATION

```
from sklearn.neighbors import KNeighborsClassifier
knn1=KNeighborsClassifier()
param={'n_neighbors':[3,5,7,9], 'weights':['uniform', 'distance']}
```

## GRIDSEARCHCV

```
from sklearn.model_selection import GridSearchCV
clf=GridSearchCV(knn1,param,cv=10,scoring='accuracy')
clf.fit(x_train,y_train)
print(clf.best_params_)

{'n_neighbors': 3, 'weights': 'distance'}

#Model creation
from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=3,weights='distance')
knn.fit(x_train,y_train)
y_prediction=knn.predict(x_test)
y_prediction
array(['yes', 'no', 'no', ..., 'no', 'yes', 'no'], dtype=object)
```

## PERFORMANCE EVALUATION

```
from sklearn.metrics import
accuracy_score,confusion_matrix,classification_report,ConfusionMatrixD
isplay
#confusion matrix
matr=confusion_matrix(y_test,y_prediction)
print(matr)

[[7469  513]
 [ 363 7611]]

#accuracy score
score=accuracy_score(y_test,y_prediction)
score
```

0.9450990223113562

*#classification report*


```
report=classification_report(y_test,y_prediction)
print(report)
```

	precision	recall	f1-score	support
no	0.95	0.94	0.94	7982
yes	0.94	0.95	0.95	7974
accuracy			0.95	15956
macro avg	0.95	0.95	0.95	15956
weighted avg	0.95	0.95	0.95	15956

*#confusion matrix display*

```
labels=['no','yes']
cmd=ConfusionMatrixDisplay(matr,display_labels=labels)
cmd.plot()
```

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay at 0x7bb7b4270580>



```

from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import BernoulliNB
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import
confusion_matrix,accuracy_score,classification_report

knn=KNeighborsClassifier(n_neighbors=7)
model2 = RandomForestClassifier(criterion =
'entropy',n_estimators=100,random_state=33)
model3 = DecisionTreeClassifier(criterion = 'entropy')
base=BernoulliNB()
model=SVC()
lst=[knn,model2,base,model,model3]

for i in lst:
    print(i)
    i.fit(x_train,y_train)
    y_pred=i.predict(x_test)
    print(confusion_matrix(y_test,y_pred))
    print(accuracy_score(y_test,y_pred))
    print(classification_report(y_test,y_pred))
    print('*'*100)

```

```

KNeighborsClassifier(n_neighbors=7)
[[7435  547]
 [ 464 7510]]
0.936638255201805

```

	precision	recall	f1-score	support
no	0.94	0.93	0.94	7982
yes	0.93	0.94	0.94	7974
accuracy			0.94	15956
macro avg	0.94	0.94	0.94	15956
weighted avg	0.94	0.94	0.94	15956

```

*****
*****

```

```

RandomForestClassifier(criterion='entropy', random_state=33)
[[7723  259]
 [ 340 7634]]
0.9624592629731762

```

	precision	recall	f1-score	support
no	0.96	0.97	0.96	7982
yes	0.97	0.96	0.96	7974
accuracy			0.96	15956



macro avg	0.96	0.96	0.96	15956
weighted avg	0.96	0.96	0.96	15956

\*\*\*\*\*  
\*\*\*\*\*

BernoulliNB()

[[7210 772]

[ 869 7105]]

0.8971546753572324

	precision	recall	f1-score	support
no	0.89	0.90	0.90	7982
yes	0.90	0.89	0.90	7974

accuracy			0.90	15956
macro avg	0.90	0.90	0.90	15956
weighted avg	0.90	0.90	0.90	15956

\*\*\*\*\*  
\*\*\*\*\*

SVC()

[[7578 404]

[ 650 7324]]

0.9339433441965405

	precision	recall	f1-score	support
no	0.92	0.95	0.93	7982
yes	0.95	0.92	0.93	7974

accuracy			0.93	15956
macro avg	0.93	0.93	0.93	15956
weighted avg	0.93	0.93	0.93	15956

\*\*\*\*\*  
\*\*\*\*\*

DecisionTreeClassifier(criterion='entropy')

[[7364 618]

[ 476 7498]]

0.9314364502381549

	precision	recall	f1-score	support
no	0.94	0.92	0.93	7982
yes	0.92	0.94	0.93	7974

accuracy			0.93	15956
macro avg	0.93	0.93	0.93	15956
weighted avg	0.93	0.93	0.93	15956

\*\*\*\*\*  
\*\*\*\*\*